

ChonkSort

Generated by Doxygen 1.8.17

1 BMSRS	1
2 Namespace Index	3
2.1 Namespace List	3
3 Namespace Documentation	5
3.1 Arrays::anonymous_namespace{sort.cpp} Namespace Reference	5
3.1.1 Function Documentation	5
3.1.1.1 bitScanRev()	6
3.1.1.2 hSort()	6
3.1.1.3 iSort()	7
3.1.1.4 parallelPrefixFill()	7
3.1.1.5 qSort()	8
3.1.1.6 siftDown()	9
3.1.2 Variable Documentation	10
3.1.2.1 DeBruijn64	10
3.1.2.2 DeBruijnTableF	10
Index	11

Chapter 1

BMSRS

Chapter 2

Namespace Index

2.1 Namespace List

Here is a list of all documented namespaces with brief descriptions:

Arrays::anonymous_namespace{sort.cpp}	5
---	---

Chapter 3

Namespace Documentation

3.1 Arrays::anonymous_namespace{sort.cpp} Namespace Reference

Functions

- template<typename E >
void [parallelPrefixFill](#) (E &x)
- template<> void **parallelPrefixFill**< [uint64_t](#) > (uint64_t &x)
- template<> void **parallelPrefixFill**< [uint32_t](#) > (uint32_t &x)
- template<> void **parallelPrefixFill**< [uint16_t](#) > (uint16_t &x)
- template<> void **parallelPrefixFill**< [uint8_t](#) > (uint8_t &x)
- template<typename E >
constexpr int [bitScanRev](#) (E l)
- template<typename E >
constexpr void **swap** (E *const i, E *const j)
- template<typename E, bool Leftmost>
void [iSort](#) (E *const low, E *const high)
- template<typename E >
void [siftDown](#) (E *const a, const int i, const int size)
- template<typename E >
void [hSort](#) (E *const low, E *const high)
- template<typename E, bool Leftmost = true>
void [qSort](#) (E *const low, E *const high, const int height)

Variables

- constexpr [uint32_t](#) **DPQSInsertionThreshold** = 32
- constexpr [uint64_t](#) [DeBruijn64](#)
- constexpr [uint8_t](#) **DeBruijnTableF** []

3.1.1 Function Documentation

3.1.1.1 bitScanRev()

```
template<typename E >
constexpr int Arrays::anonymous_namespace{sort.cpp}::bitScanRev (
    E l ) [constexpr]
```

bitScanReverse

Authors

Kim Walisch
Mark Dickinson

Parameters

<i>bb</i>	bitboard to scan @precondition <i>bb</i> != 0
-----------	---

Returns

index (0..63) of most significant one bit

3.1.1.2 hSort()

```
template<typename E >
void Arrays::anonymous_namespace{sort.cpp}::hSort (
    E *const low,
    E *const high ) [inline]
```

Heap Sort

Classical heap sort that sorts the given range in ascending order, building a max heap and continuously sifting/swapping the max element to the previous rightmost index.

Author

Ellie Moore

Template Parameters

<i>E</i>	the element type
----------	------------------

Parameters

<i>low</i>	a pointer to the leftmost index
<i>high</i>	a pointer to the rightmost index

3.1.1.3 iSort()

```
template<typename E , bool Leftmost>
void Arrays::anonymous_namespace{sort.cpp}::iSort (
    E *const low,
    E *const high ) [inline]
```

Insertion Sort

Classical ascending insertion sort packaged with a "pairing" optimization to be used in the context of Quick Sort. This optimization is used whenever the portion of the array to be sorted is padded on the left by a portion with lesser elements. The fact that all of the elements on the left are automatically less than the elements in the current portion allows us to skip the costly lower boundary check in the nested loops.

Authors

Josh Blosh

John Bently

Ellie Moore

Template Parameters

<i>E</i>	the element type
----------	------------------

Parameters

<i>low</i>	a pointer to the leftmost index
<i>high</i>	a pointer to the rightmost index
<i>isLeftmost</i>	whether or not this subarray is a left-most partition.

3.1.1.4 parallelPrefixFill()

```
template<typename E >
void Arrays::anonymous_namespace{sort.cpp}::parallelPrefixFill (
    E & x )
```

Fill trailing bits using prefix fill.

Example:

```
10000000 » 1
= 01000000 | 10000000
= 11000000 » 2
= 00110000 | 11000000
= 11110000 » 4
= 00001111 | 11110000
= 11111111
```

Template Parameters

<i>E</i>	The type
----------	----------

Parameters

<code>x</code>	The integer
----------------	-------------

3.1.1.5 qSort()

```
template<typename E , bool Leftmost = true>
void Arrays::anonymous_namespace{sort.cpp}::qSort (
    E *const low,
    E *const high,
    const int height )
```

Introspective Multi-Pivot Quick Sort

This sort uses a recursive pattern to re-arrange data in ascending order. Whenever the size of the current interval falls below the threshold of 32 elements, insertion sort is used. On return, the data in the current interval will be in sorted order.

<https://codeblab.com/wp-content/uploads/2009/09/DualPivotQuicksort.pdf>

Here is the original Dual-Pivot Quicksort by Vladimir Yaroslavskiy.

https://www.researchgate.net/publication/289974363_Multi-Pivot_Quicksort_Theory_and_Experiments

Here is the three pivot partition algorithm by Kushagra, Lopez-Ortiz, Munro, and Qiao.

An optimized Dual-Pivot Quicksort by Bloch and Bently can be found in the open-source jdk within the java.util library.

This sort is derived from the resources listed above.

Optimizations:**1. Use Smart Pivot Selection**

Pivots are carefully selected from a set of five candidates. These candidates are the leftmost element, the element at the first tercile, the middle element, the element at the second tercile, and the rightmost element. All five of these candidates are sorted in place. If none are equal, three pivots are chosen from the middle. If an outside pair is equal, two pivots are selected from the tercile indices. However, if a middle pair is equal, then a single pivot will be selected from the middle and traditional Quick Sort will be used. This process helps to select pivots that divide the data as evenly as possible, mitigating the possibility of the worst-case $O(N^2)$ time complexity.

2. Use Four-Way or Three-Way Partitioning When Possible

Multi-way partitioning results in smaller intervals to partition in recursive calls. This makes cache hits more likely. Multi-way partitioning also helps to trim the height of the sort tree.

3. Ignore In-Order and Equal Elements

Runs of elements will be ignored.

4. Keep the Middle Portion(s) as Small as Possible

After partitioning with more than one pivot, if the middle portion comprises more than 2/3 of the current interval, we swap all elements equal to the pivots out of the way before recursively sorting.

5. Use Insertion Sort On Little Sub-Arrays

Insertion Sort is small and relies heavily on swapping. Fewer instructions help to alleviate the overhead that plagues sorting algorithms, and consistent swapping leverages CPU caching to its fullest extent. These features allow insertion sort to beat the runtimes of $O(n \log n)$ sorts on small sets of data.

6. Use Heap Sort to prevent Quicksort's worst-case time complexity.

This sort is an introspective sort&mdash a sort that switches to a guaranteed $O(n \log n)$ algorithm whenever time complexity is becoming quadratic. Here, Heap Sort is used to avoid the overhead of auxilliary storage.

Authors

Vladimir Yaroslavskiy
 Josh Bloch
 Jon Bently
 Shrinu Kushagra
 Alejandro Lopez-Ortiz
 J. Ian Munro
 Aurick Qiao
 Ellie Moore

Template Parameters

<i>E</i>	the element type
<i>Leftmost</i>	whether we are considering a leftmost partition.

Parameters

<i>low</i>	a pointer to the leftmost index
<i>high</i>	a pointer to the rightmost index
<i>Leftmost</i>	whether or not this subarray is a left-most partition.

3.1.1.6 siftDown()

```
template<typename E >
void Arrays::anonymous_namespace{sort.cpp}::siftDown (
    E *const a,
    const int i,
    const int size ) [inline]
```

A generic "sift down" method (AKA max-heapify)

Template Parameters

<i>E</i>	the element type
----------	------------------

Parameters

<i>a</i>	the pointer to the base of the current sub-array
<i>i</i>	the starting index
<i>size</i>	the size of the current sub-array

3.1.2 Variable Documentation

3.1.2.1 DeBruijn64

```
constexpr uint64_t Arrays::anonymous_namespace{sort.cpp}::DeBruijn64 [constexpr]
```

Initial value:

```
=
    0x03F79D71B4CB0A89L
```

The DeBruijn constant.

3.1.2.2 DeBruijnTableF

```
constexpr uint8_t Arrays::anonymous_namespace{sort.cpp}::DeBruijnTableF[] [constexpr]
```

Initial value:

```
= {
    0,  47,  1, 56, 48, 27,  2, 60,
   57, 49, 41, 37, 28, 16,  3, 61,
   54, 58, 35, 52, 50, 42, 21, 44,
   38, 32, 29, 23, 17, 11,  4, 62,
   46, 55, 26, 59, 40, 36, 15, 53,
   34, 51, 20, 43, 31, 22, 10, 45,
   25, 39, 14, 33, 19, 30,  9, 24,
   13, 18,  8, 12,  7,  6,  5, 63
}
```

Index

Arrays::anonymous_namespace{sort.cpp}, [5](#)
 bitScanRev, [5](#)
 DeBruijn64, [10](#)
 DeBruijnTableF, [10](#)
 hSort, [6](#)
 iSort, [6](#)
 parallelPrefixFill, [7](#)
 qSort, [8](#)
 siftDown, [9](#)

bitScanRev
 Arrays::anonymous_namespace{sort.cpp}, [5](#)

DeBruijn64
 Arrays::anonymous_namespace{sort.cpp}, [10](#)

DeBruijnTableF
 Arrays::anonymous_namespace{sort.cpp}, [10](#)

hSort
 Arrays::anonymous_namespace{sort.cpp}, [6](#)

iSort
 Arrays::anonymous_namespace{sort.cpp}, [6](#)

parallelPrefixFill
 Arrays::anonymous_namespace{sort.cpp}, [7](#)

qSort
 Arrays::anonymous_namespace{sort.cpp}, [8](#)

siftDown
 Arrays::anonymous_namespace{sort.cpp}, [9](#)